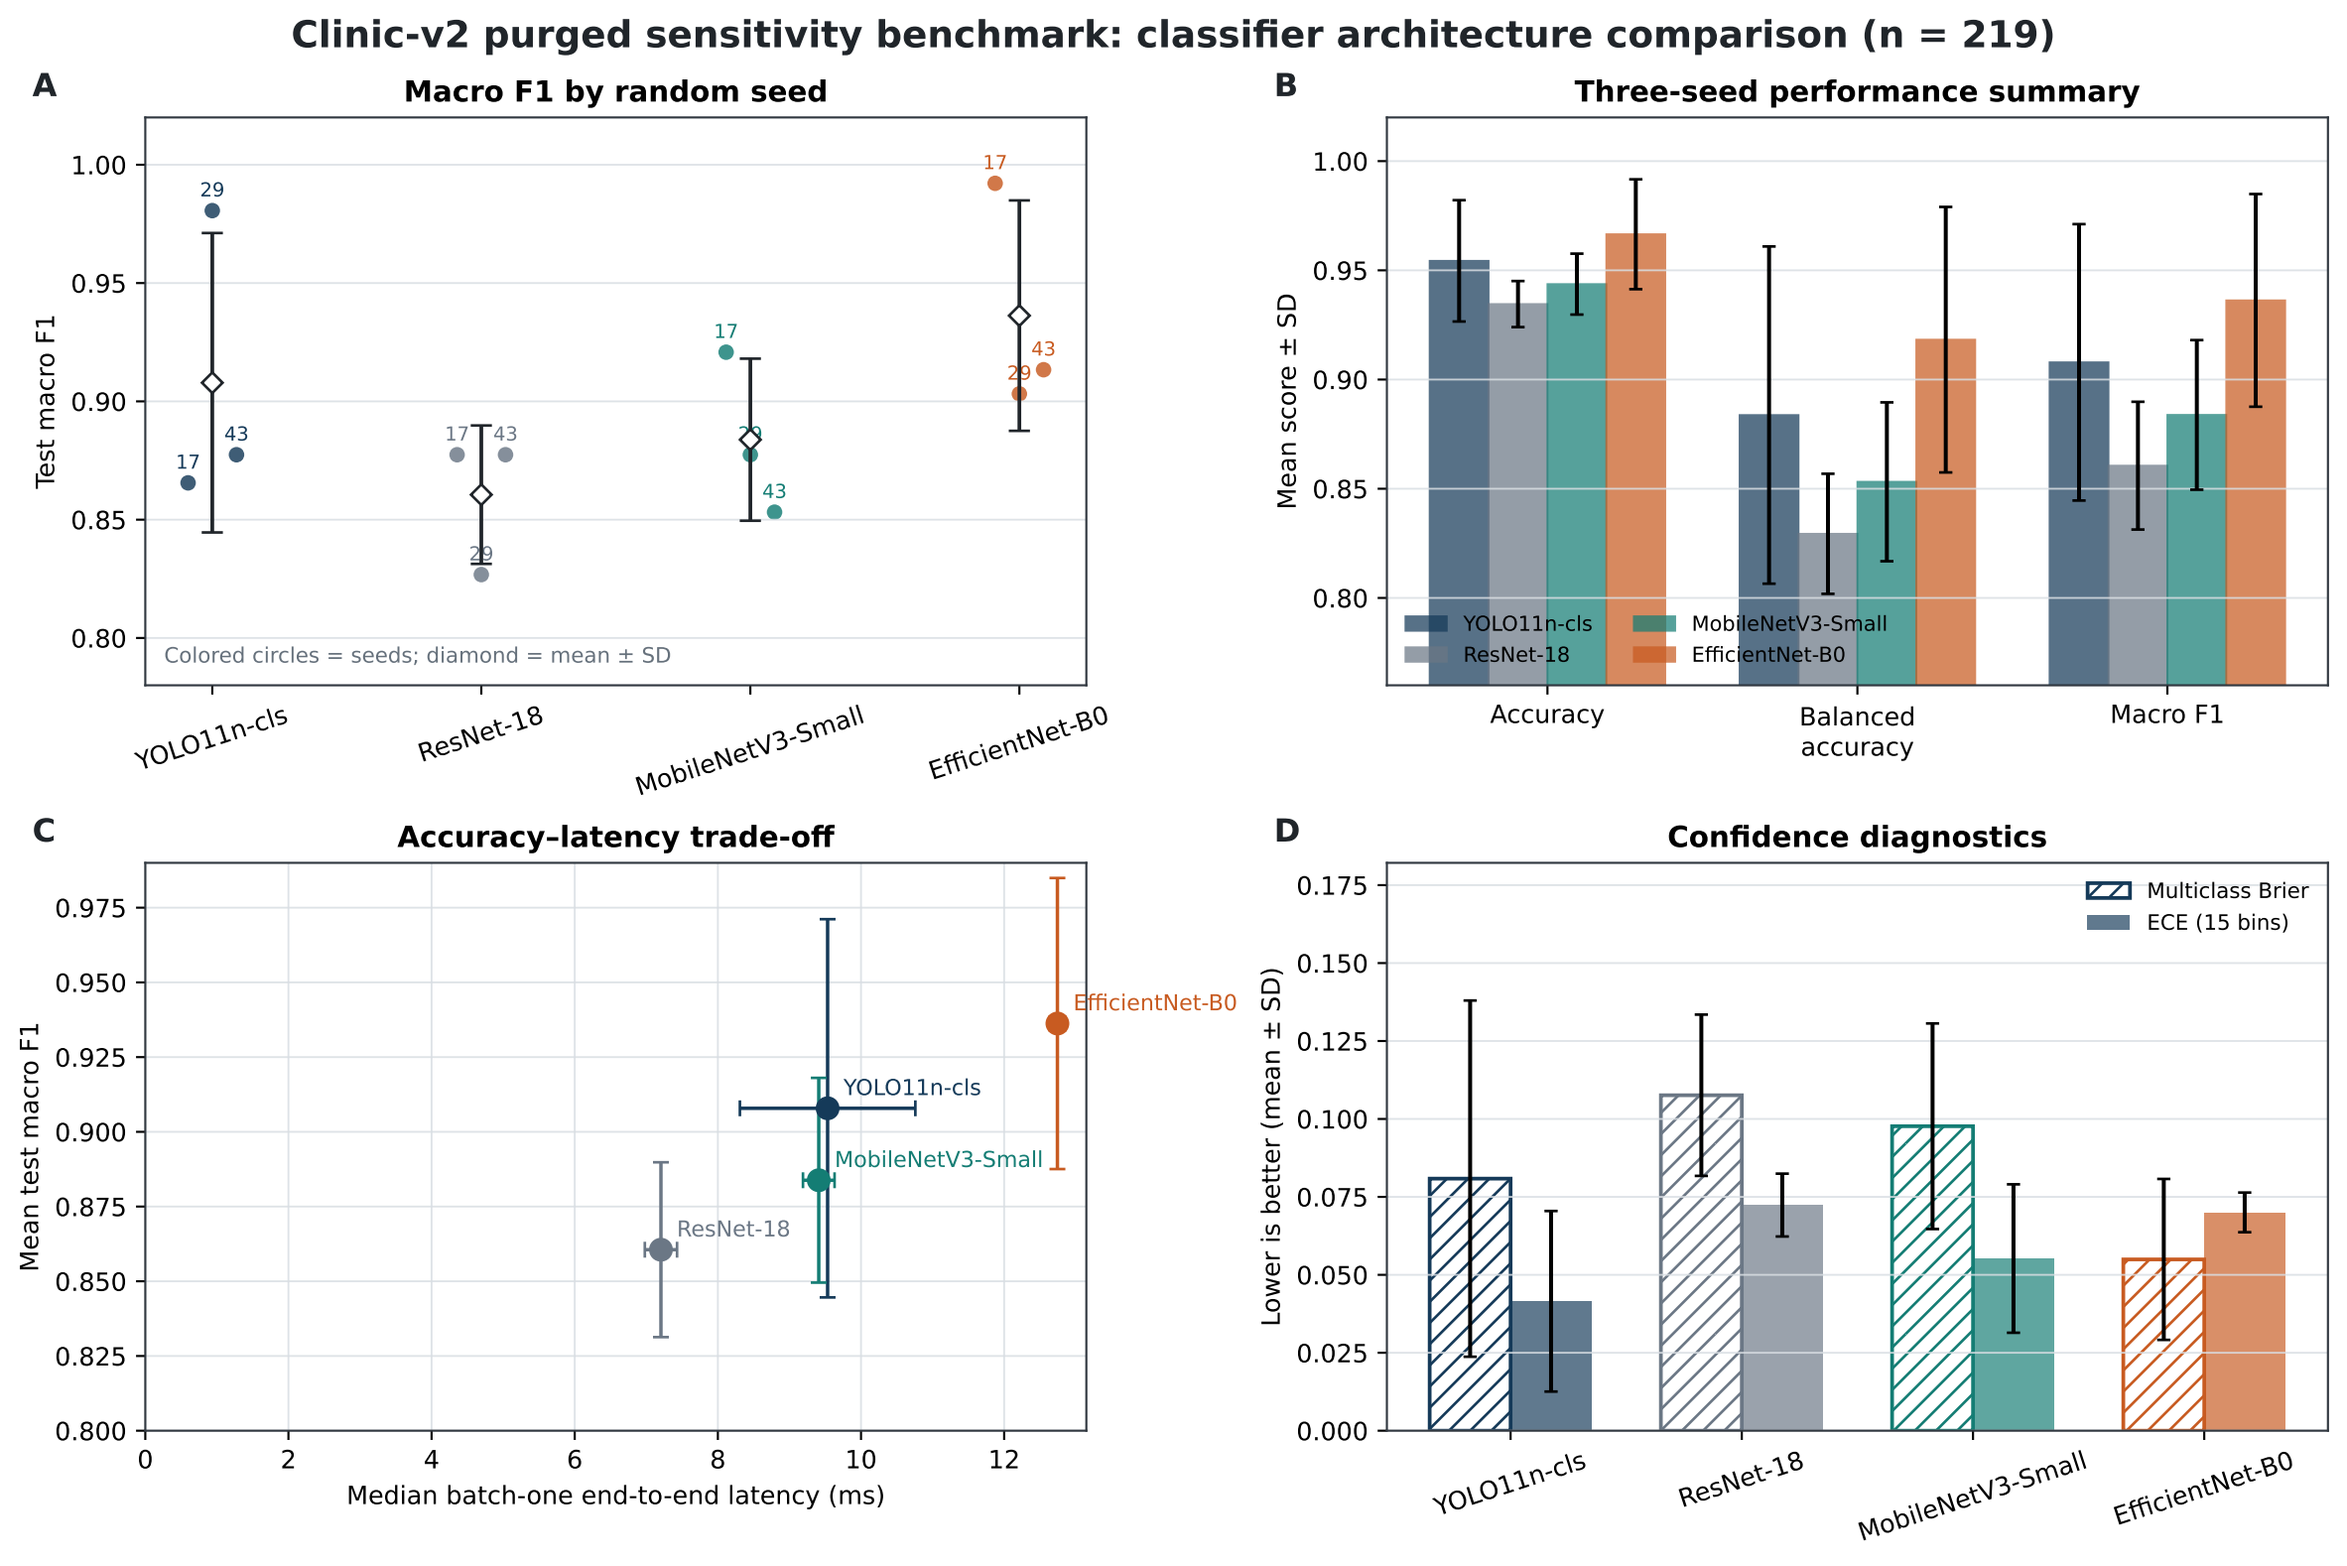
Clinic-v2 purged sensitivity benchmark: classifier architecture comparison (n = 219)
A
Macro F1 by random seed
Test macro F1
1.00
0.95
0.90
0.85
0.80
Colored circles = seeds; diamond = mean ± SD
YOLO11n-cls
ResNet-18
MobileNetV3-Small
EfficientNet-B0
B
Three-seed performance summary
Mean score ± SD
YOLO11n-cls
ResNet-18
MobileNetV3-Small
EfficientNet-B0
Accuracy
Balanced accuracy
Macro F1
C
Accuracy–latency trade-off
Mean test macro F1
0.975
0.950
0.925
0.900
0.875
0.850
0.825
0.800
0
2
4
6
8
10
12
Median batch-one end-to-end latency (ms)
YOLO11n-cls
MobileNetV3-Small
ResNet-18
EfficientNet-B0
D
Confidence diagnostics
Lower is better (mean ± SD)
0.175
0.150
0.125
0.100
0.075
0.050
0.025
0.000
Multiclass Brier
ECE (15 bins)
YOLO11n-cls
ResNet-18
MobileNetV3-Small
EfficientNet-B0