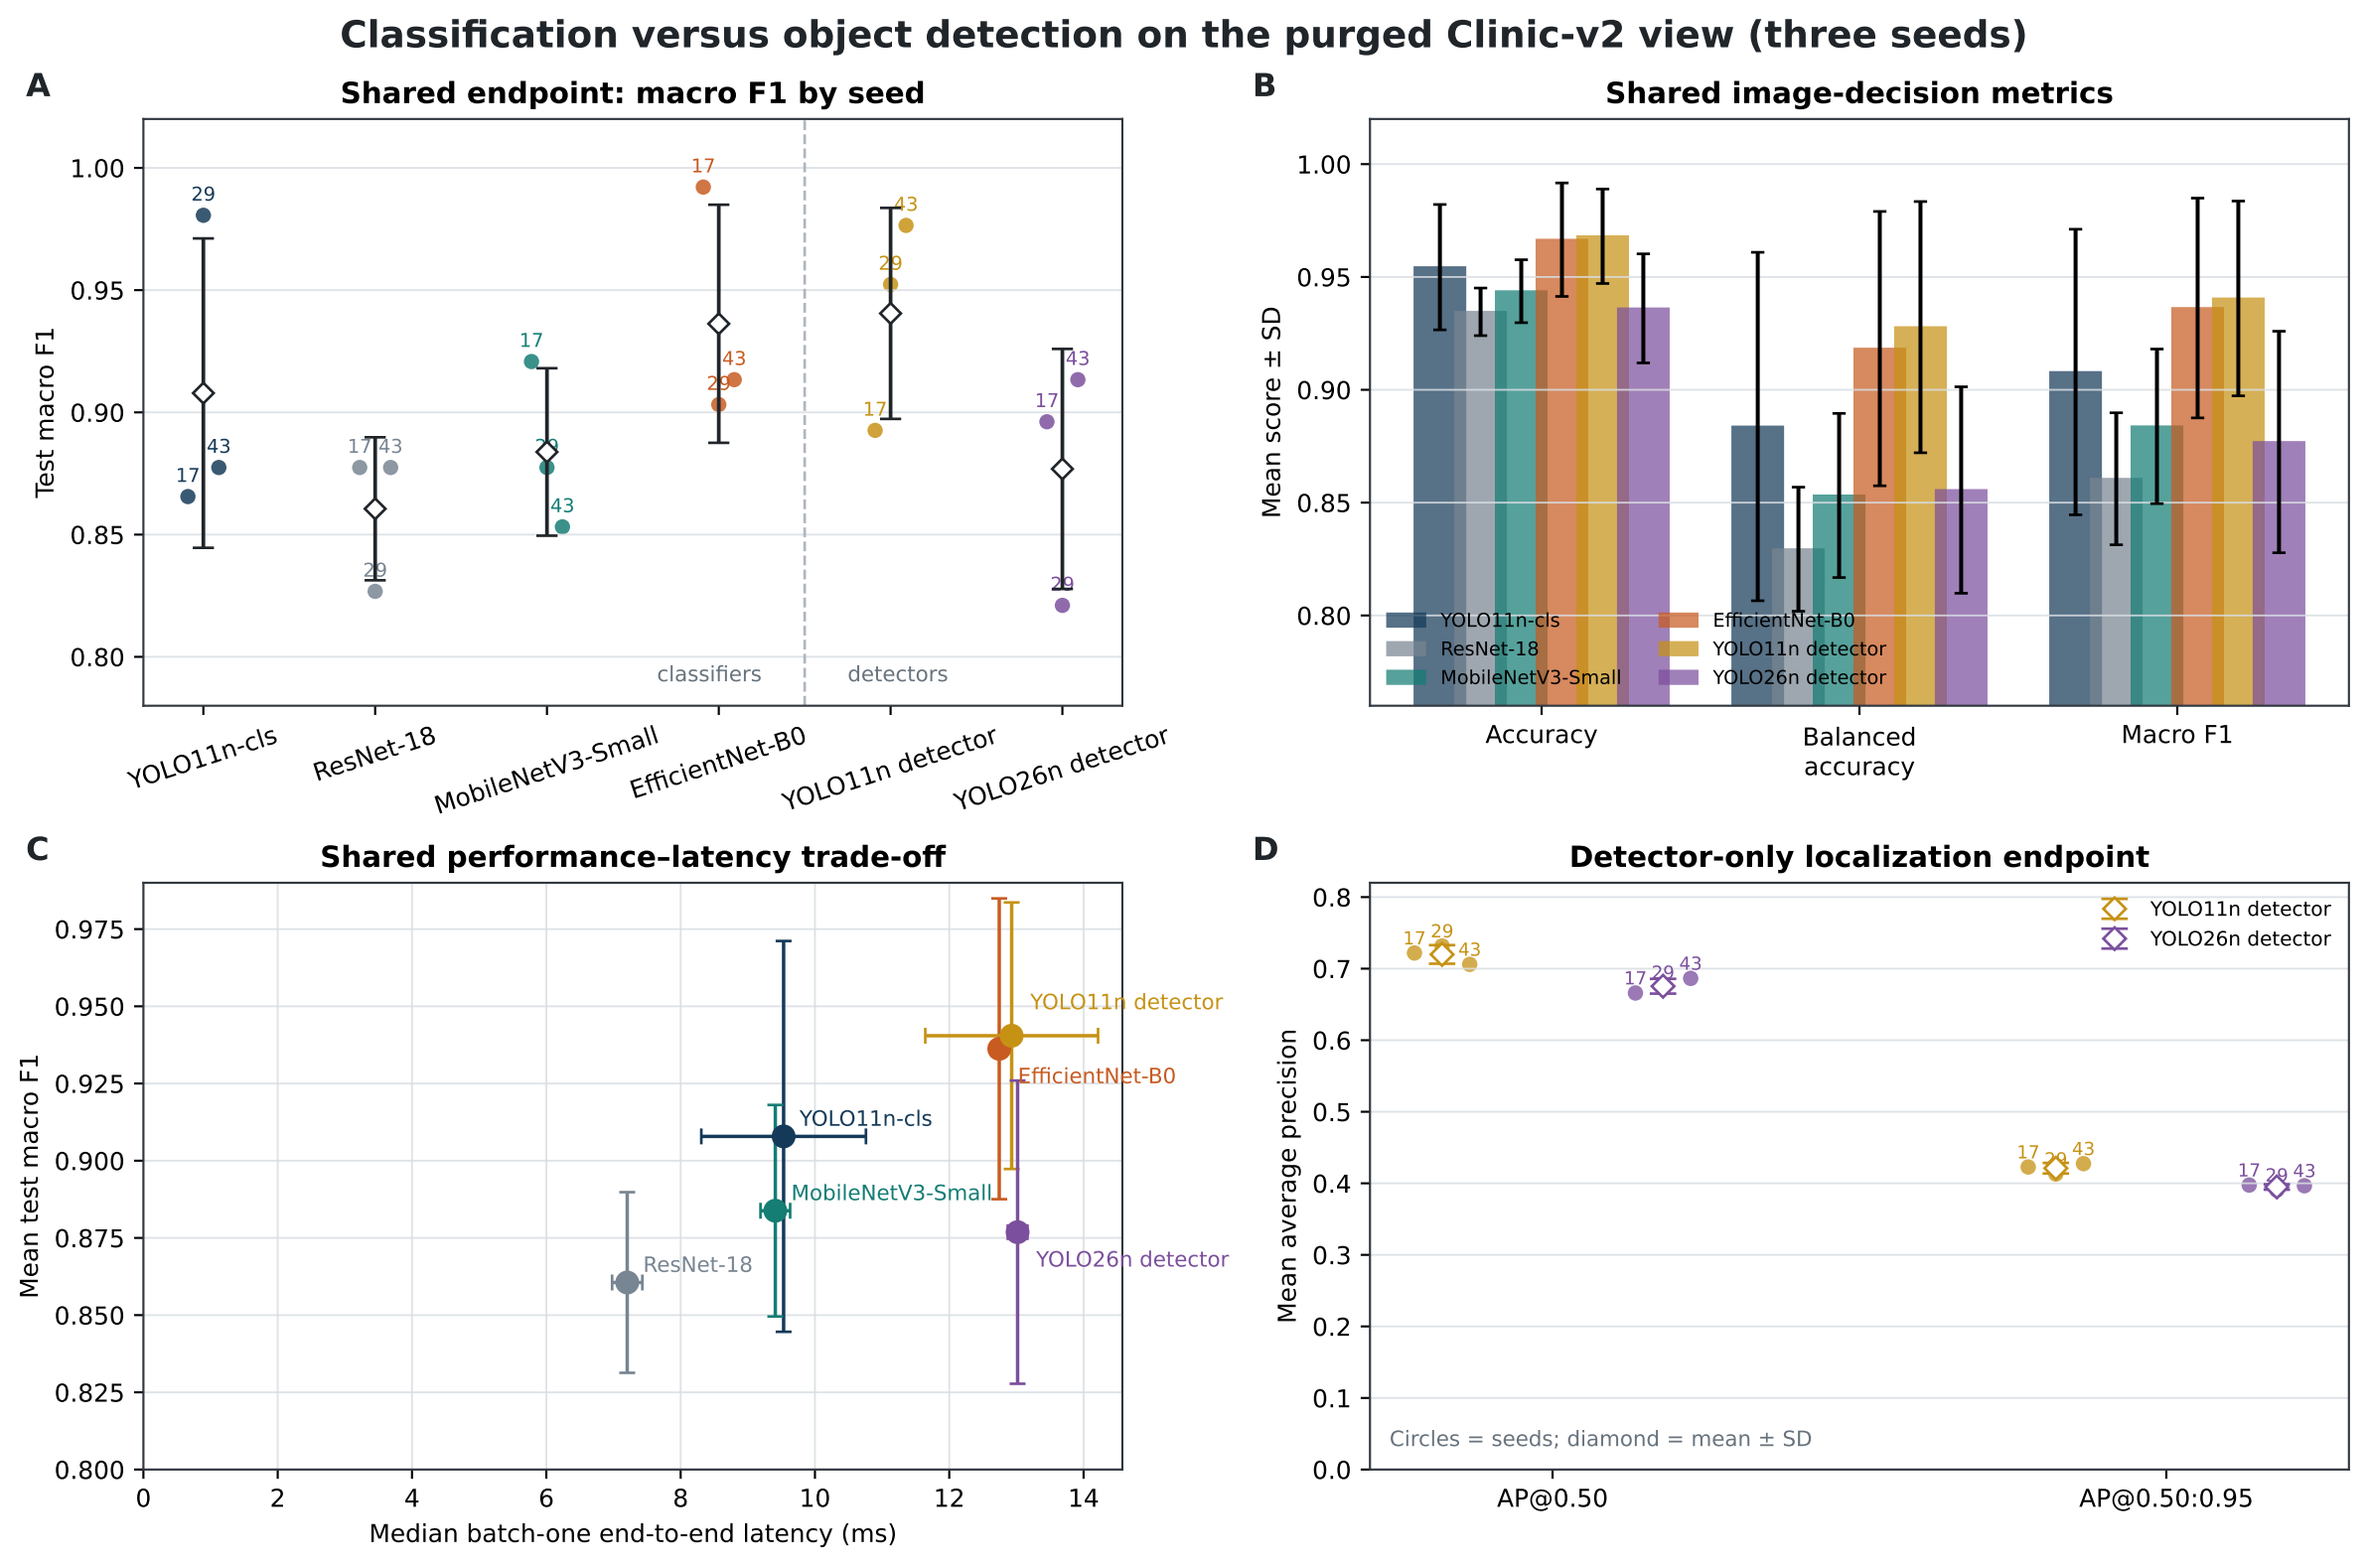
Classification versus object detection on the purged Clinic-v2 view (three seeds)
A
Shared endpoint: macro F1 by seed
Test macro F1
1.00
0.95
0.90
0.85
0.80
classifiers
detectors
YOLO11n-cls
ResNet-18
MobileNetV3-Small
EfficientNet-B0
YOLO11n detector
YOLO26n detector
B
Shared image-decision metrics
Mean score ± SD
YOLO11n-cls
ResNet-18
MobileNetV3-Small
EfficientNet-B0
YOLO11n detector
YOLO26n detector
Accuracy
Balanced accuracy
Macro F1
C
Shared performance–latency trade-off
Mean test macro F1
Median batch-one end-to-end latency (ms)
YOLO11n detector
EfficientNet-B0
YOLO11n-cls
MobileNetV3-Small
ResNet-18
YOLO26n detector
D
Detector-only localization endpoint
Mean average precision
YOLO11n detector
YOLO26n detector
Circles = seeds; diamond = mean ± SD
AP@0.50
AP@0.50:0.95